

Class 8

Challana Tea

```
#Import the dataset
```

Before we can begin our analysis we first have to download and import the data correctly into our R session. Use `read.csv()`

```
wisc.df <- read.csv("Wisconsincancer.csv",row.names=1)
head(wisc.df)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.10	1297.0
843786	M	12.45	15.70	82.57	477.1
	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302	0.11840	0.27760	0.3001	0.14710	
842517	0.08474	0.07864	0.0869	0.07017	
84300903	0.10960	0.15990	0.1974	0.12790	
84348301	0.14250	0.28390	0.2414	0.10520	
84358402	0.10030	0.13280	0.1980	0.10430	
843786	0.12780	0.17000	0.1578	0.08089	
	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419	0.07871	1.0950	0.9053	8.589
842517	0.1812	0.05667	0.5435	0.7339	3.398
84300903	0.2069	0.05999	0.7456	0.7869	4.585
84348301	0.2597	0.09744	0.4956	1.1560	3.445
84358402	0.1809	0.05883	0.7572	0.7813	5.438
843786	0.2087	0.07613	0.3345	0.8902	2.217
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
842302	153.40	0.006399	0.04904	0.05373	0.01587
842517	74.08	0.005225	0.01308	0.01860	0.01340

84300903	94.03	0.006150	0.04006	0.03832	0.02058
84348301	27.23	0.009110	0.07458	0.05661	0.01867
84358402	94.44	0.011490	0.02461	0.05688	0.01885
843786	27.19	0.007510	0.03345	0.03672	0.01137
		symmetry_se fractal_dimension_se radius_worst texture_worst			
842302	0.03003	0.006193	25.38	17.33	
842517	0.01389	0.003532	24.99	23.41	
84300903	0.02250	0.004571	23.57	25.53	
84348301	0.05963	0.009208	14.91	26.50	
84358402	0.01756	0.005115	22.54	16.67	
843786	0.02165	0.005082	15.47	23.75	
		perimeter_worst area_worst smoothness_worst compactness_worst			
842302	184.60	2019.0	0.1622	0.6656	
842517	158.80	1956.0	0.1238	0.1866	
84300903	152.50	1709.0	0.1444	0.4245	
84348301	98.87	567.7	0.2098	0.8663	
84358402	152.20	1575.0	0.1374	0.2050	
843786	103.40	741.6	0.1791	0.5249	
		concavity_worst concave.points_worst symmetry_worst			
842302	0.7119	0.2654	0.4601		
842517	0.2416	0.1860	0.2750		
84300903	0.4504	0.2430	0.3613		
84348301	0.6869	0.2575	0.6638		
84358402	0.4000	0.1625	0.2364		
843786	0.5355	0.1741	0.3985		
		fractal_dimension_worst X			
842302		0.11890 NA			
842517		0.08902 NA			
84300903		0.08758 NA			
84348301		0.17300 NA			
84358402		0.07678 NA			
843786		0.12440 NA			

Q1. How many samples and variables are in this dataset?

```
nrow(wisc.df)
```

```
[1] 569
```

```
ncol(wisc.df)
```

```
[1] 32
```

Q2. How many M and B samples are there?

```
sum(wisc.df$diagnosis=="M")
```

```
[1] 212
```

```
sum(wisc.df$diagnosis=="B")
```

```
[1] 357
```

```
table(wisc.df$diagnosis)
```

	B	M
357	212	

Q3. How many variables/features in the data are suffixed with _mean?

```
length(grep("_mean", colnames(wisc.df)))
```

```
[1] 10
```

Q. What features are “mean” values (or have “mean” in their name)

```
(grep("_mean", colnames(wisc.df), value=TRUE))
```

```
[1] "radius_mean"           "texture_mean"          "perimeter_mean"  
[4] "area_mean"             "smoothness_mean"       "compactness_mean"  
[7] "concavity_mean"        "concave.points_mean"   "symmetry_mean"  
[10] "fractal_dimension_mean"
```

I need to remove the first diagnosis column from the data before doing any analysis. I will store it for later as a factor

```
# We can use -1 here to remove the first column
wisc.data <- wisc.df[,-1]
diagnosis <- as.factor(wisc.df$diagnosis)
```

2. Principal Component Analysis

The main PCA function in base R is called `prcom()`.

Before doing anything like PCA, it is important to check if the data need to be scaled before performing PCA. Recall two common reasons for scaling data include:

- The input variables use different units of measurement.
- The input variables have significantly different variances.

Looks like we need to scale by setting `scale=TRUE` in our `prcomp()` function call.

```
round(apply(wisc.data, 2, sd), 2)
```

	radius_mean	texture_mean	perimeter_mean
	3.52	4.30	24.30
	area_mean	smoothness_mean	compactness_mean
	351.91	0.01	0.05
	concavity_mean	concave.points_mean	symmetry_mean
	0.08	0.04	0.03
fractal_dimension_mean		radius_se	texture_se
	0.01	0.28	0.55
perimeter_se		area_se	smoothness_se
	2.02	45.49	0.00
compactness_se		concavity_se	concave.points_se
	0.02	0.03	0.01
symmetry_se		fractal_dimension_se	radius_worst
	0.01	0.00	4.83
texture_worst		perimeter_worst	area_worst
	6.15	33.60	569.36
smoothness_worst		compactness_worst	concavity_worst
	0.02	0.16	0.21
concave.points_worst		symmetry_worst	fractal_dimension_worst
	0.07	0.06	0.02
X			
NA			

```
wisc.data <- wisc.df[,-32]

wisc.data <- wisc.data[,-1]
```

Time for PCA

```
wisc.pr <- prcomp(wisc.data, scale=TRUE)
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

44.27

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

PC3

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

PC7

```
attributes(wisc.pr)

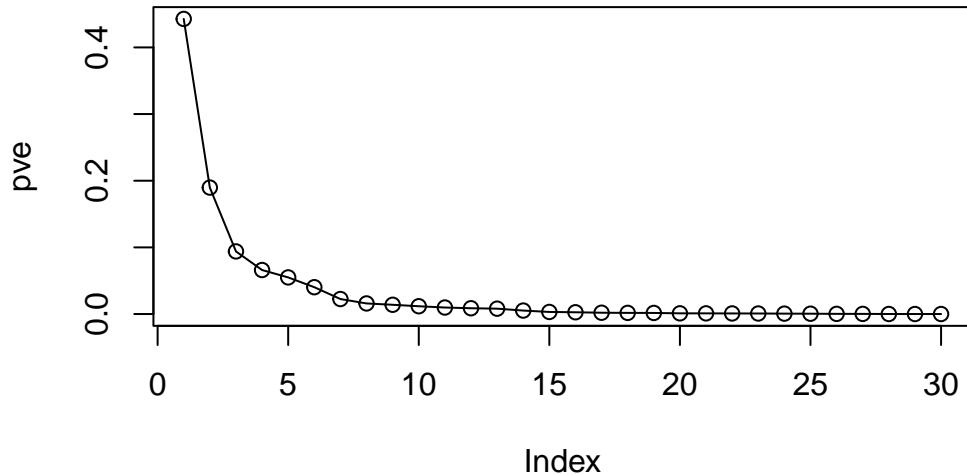
$names
[1] "sdev"      "rotation"   "center"     "scale"      "x"

$class
[1] "prcomp"
```

Make a little scree plot

```
pr.var <- wisc.pr$sdev^2
#Proportion of variance
pve <- pr.var/sum(pr.var)

plot(pve, typ="o")
```



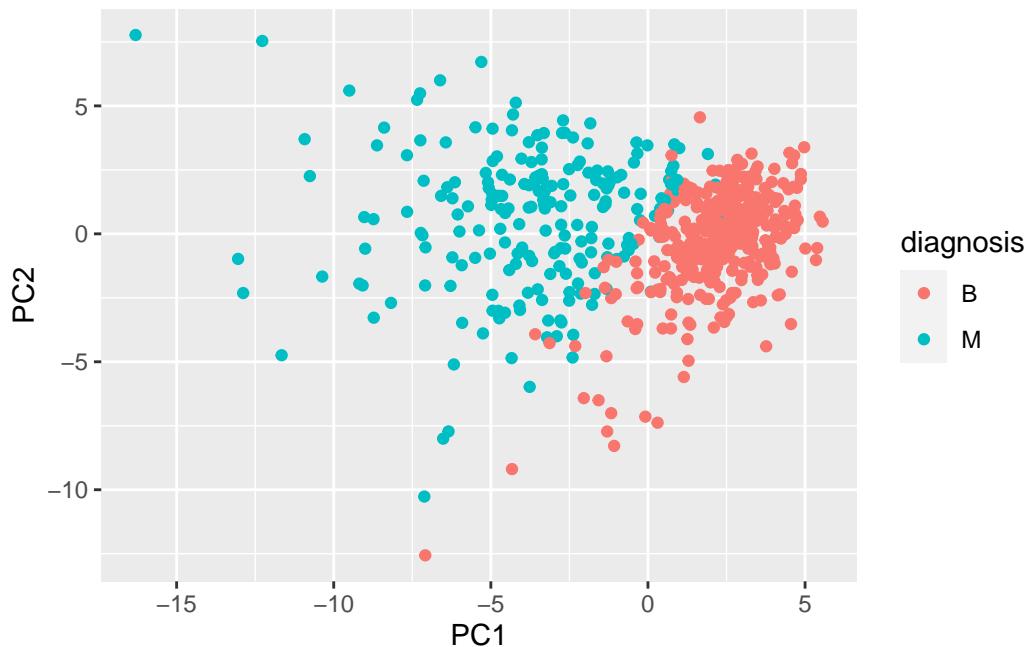
PC plot

```

library(ggplot2)
pc <- as.data.frame(wisc.pr$x)

ggplot(pc) +
  aes(x=PC1,y=PC2,col=diagnosis) +
  geom_point()

```



Hierarchical clustering

We can try clustering the original data with `hclust()` or `kmeans()`. First we will scale the data just like we did for PCA

```

data.scaled <- scale(wisc.data)
head(apply(data.scaled,2,sd))

```

	radius_mean	texture_mean	perimeter_mean	area_mean
smoothness_mean	1	1	1	1
compactness_mean				
	1	1		

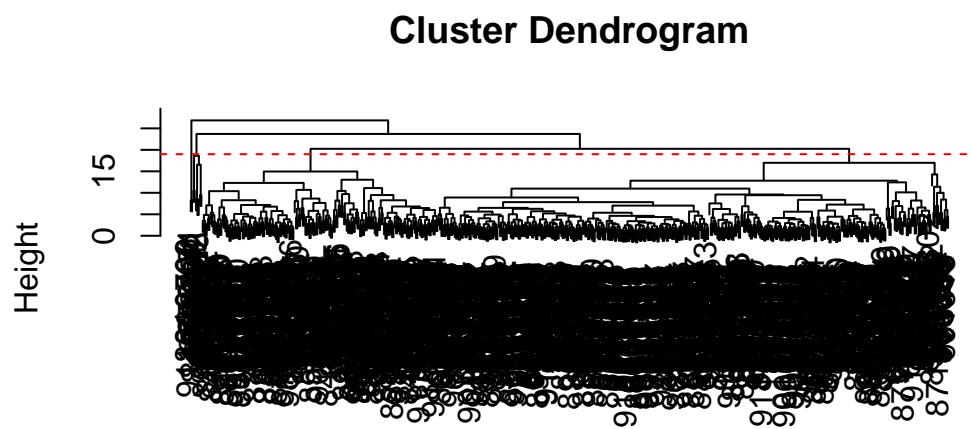
```
data.dist <- dist(data.scaled)
wisc.hclust <- hclust(data.dist, method = "complete")
wisc.hclust
```

Call:
hclust(d = data.dist, method = "complete")

Cluster method : complete
Distance : euclidean
Number of objects: 569

Plot the dendrogram

```
plot(wisc.hclust)
abline(h=19, col="red", lty=2)
```



data.dist
hclust (*, "complete")

To get a cluster membership vector I will use the `cutree()` function and “cut” into 4 or so grps or clusters.

Q10. Using the `plot()` and `abline()` functions, what is the height at which the clustering model has 4 clusters?

```
h = 19
```

```
grps <- cutree(wisc.hclust, h=19)
table(grps)
```

```
grps
 1   2   3   4
177  7 383  2
```

I can also use the `table()` to cross tabulate...

```
table(diagnosis)
```

```
diagnosis
B   M
357 212
```

```
table(grps,diagnosis)
```

```
diagnosis
grps  B   M
 1   12 165
 2    2   5
 3 343  40
 4    0   2
```

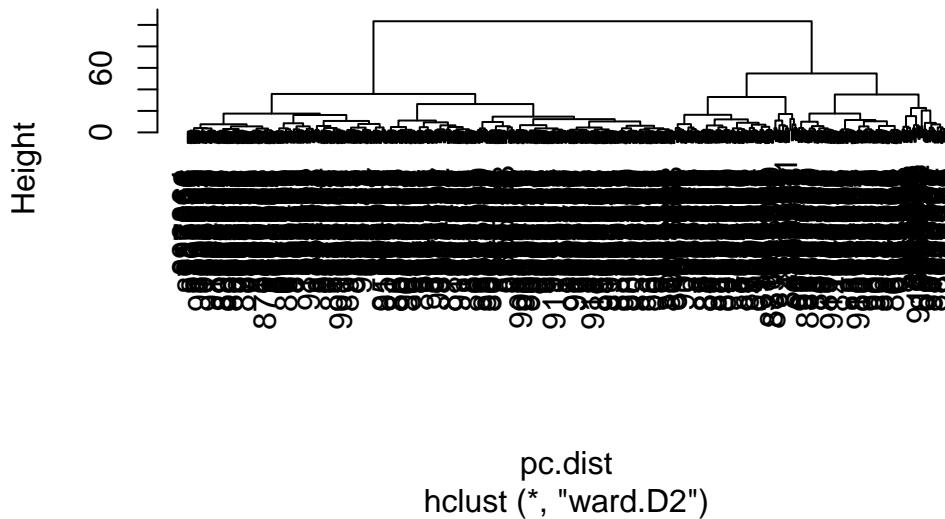
Clustering on PCA Results

I can cluster in PC-space and use as many or as few PCs as I want.

To start with I will use 3 PCs, that is I will cluster along PC1, PC2, and PC3.

```
pc.dist <- dist(wisc.pr$x[,1:3])
wisc.pr.hclust <- hclust(pc.dist, method="ward.D2")
plot(wisc.pr.hclust)
```

Cluster Dendrogram



This looks nicer than the previous clustering result. Find the two major clusters with the `cutree()` function

```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
grps
 1 2
203 366
```

```
table(grps, diagnosis)
```

grps	B	M
1	24	179
2	333	33

We could calculate accuracy - the proportion of samples correct if we take cluster 1 to represent all M and cluster 2 to represent all B.

```
(179+333)/nrow(wisc.data)
```

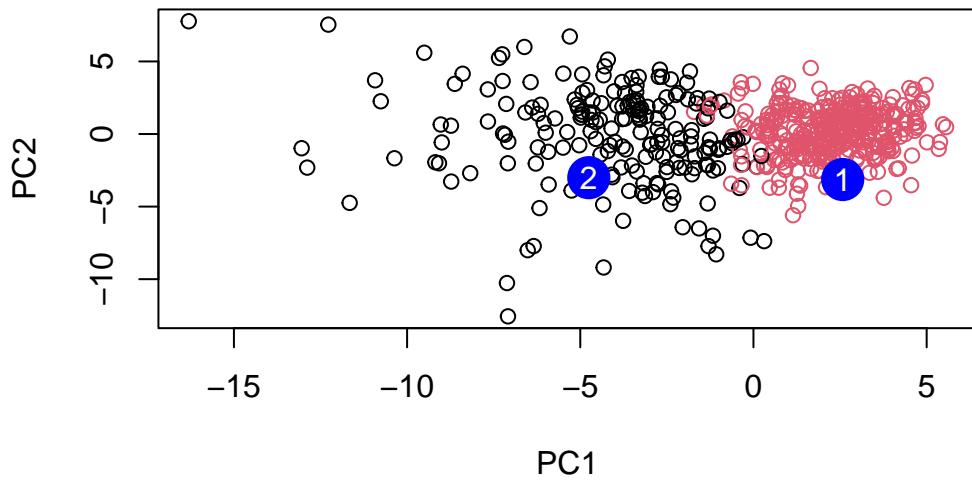
```
[1] 0.8998243
```

Prediction

```
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
[1,]	2.576616	-3.135913	1.3990492	-0.7631950	2.781648	-0.8150185	-0.3959098
[2,]	-4.754928	-3.009033	-0.1660946	-0.6052952	-1.140698	-1.2189945	0.8193031
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
[1,]	-0.2307350	0.1029569	-0.9272861	0.3411457	0.375921	0.1610764	1.187882
[2,]	-0.3307423	0.5281896	-0.4855301	0.7173233	-1.185917	0.5893856	0.303029
	PC15	PC16	PC17	PC18	PC19	PC20	
[1,]	0.3216974	-0.1743616	-0.07875393	-0.11207028	-0.08802955	-0.2495216	
[2,]	0.1299153	0.1448061	-0.40509706	0.06565549	0.25591230	-0.4289500	
	PC21	PC22	PC23	PC24	PC25	PC26	
[1,]	0.1228233	0.09358453	0.08347651	0.1223396	0.02124121	0.078884581	
[2,]	-0.1224776	0.01732146	0.06316631	-0.2338618	-0.20755948	-0.009833238	
	PC27	PC28	PC29	PC30			
[1,]	0.220199544	-0.02946023	-0.015620933	0.005269029			
[2,]	-0.001134152	0.09638361	0.002795349	-0.019015820			

```
plot(wisc.pr$x[,1:2], col=grps)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



Q16. Which of these new patients should we prioritize for follow up based on your results?

Patient 2 should be prioritized.